

# Data-centric AI: Perspectives and Challenges

Daochen Zha\*

Zaid Pervaiz Bhat†

Kwei-Herng Lai\*

Fan Yang\*

Xia Hu\*

## Abstract

The role of data in building AI systems has recently been significantly magnified by the emerging concept of data-centric AI (DCAI), which advocates a fundamental shift from model advancements to ensuring data quality and reliability. Although our community has continuously invested efforts into enhancing data in different aspects, they are often isolated initiatives on specific tasks. To facilitate the collective initiative in our community and push forward DCAI, we draw a big picture and bring together three general missions: training data development, evaluation data development, and data maintenance. We provide a top-level discussion on representative DCAI tasks and share perspectives. Finally, we list open challenges to motivate future exploration.

## 1 Introduction

Data is an indispensable element of AI systems. Recently, its role has been significantly magnified by the emerging concept of data-centric AI (DCAI), shown on the left of Figure 1. The popularity of DCAI is mainly driven by a campaign launched by Ng et al. [23], which advocates for a more data-centric rather than model-centric strategy for machine learning, a fundamental shift from model design to data quality and reliability.

The right-hand side of Figure 1 illustrates a high-level comparison between model-centric AI and DCAI. In the conventional model-centric lifecycle, the benchmark dataset remains mostly unchanged. The primary goal of the researchers and practitioners is to iterate the model to improve performance. Although this paradigm encourages model advancements, it trusts too much in data; however, data could be susceptible to undesirable flaws, which raises a question: *can the model performance reflect the actual capability, or is it just overfitting the dataset?* “Garbage in, garbage out” is often one of the first lessons we have learned in machine learning. That being said, data is not merely a fuel for AI but rather a determining factor of the model quality. DCAI has become a recent trend to shift our focus from models toward data. Attention to data in our community




Figure 1: **Left:** Tendency of DCAI over the past five years. The statistics are collected by querying Google Scholar with the exactly matched phrase “data-centric AI”. **Right:** Model-centric AI versus DCAI.

can help build more powerful AI systems to deal with more complex real-world problems.

In this paper, we define DCAI as a class of systematic techniques that develop, iterate, and maintain data for AI systems. While DCAI appears to be a new concept [26, 14, 13], many relevant research topics are not new. Our community has continuously invested efforts into enhancing data in different aspects, especially training data development. For instance, data augmentation [33] has been extensively investigated to improve data diversity by adding slightly modified data samples or new synthetic samples into training sets; feature selection [19] has been studied since decades ago for preparing cleaner and more understandable data.

Despite these individual initiatives on specific tasks, there is a lack of a top-level summary and outlook of DCAI. In particular, the role of data stretches well beyond constructing data for training. First, it is equally crucial to build a novel evaluation set to assess and understand the model quality comprehensively. Second, in industrial applications, data is not created once but rather demands continuous maintenance. It is essential to develop efficient algorithms, tools, and infrastructures to organize, understand, and debug data.

To facilitate the collective initiative in our community and push forward DCAI, we draw a big picture and share our perspectives with peers. Guided by Figure 2, our discussion covers three general DCAI missions: training data development (Section 2), evaluation data development (Section 3), and data maintenance (Section 4). Then, we discuss open challenges in Section 5, followed by a conclusion in Section 6.

\*Rice University, {daochen.zha,khlai,fyang,Xia.Hu}@rice.edu

†Texas A&M University, zaid.bhat1234@tamu.edu




Figure 2: A big picture of DCAI missions and their representative tasks/subtasks.

## 2 Training Data Development

Training data is a collection of data instances used to teach machine learning models. Constructing high-quality training data is critical to achieving DCAI. In general, there are many tasks in the data construction process, which can be divided into five high-level goals: 1) data collection, 2) data labeling, 3) data preparation, 4) data reduction, and 5) data augmentation.

**Data collection** is the process of gathering data. A straightforward approach is to construct new datasets from scratch. However, this is often time-consuming. Thus, more efficient methods are proposed by leveraging the existing datasets. A representative task is *dataset discovery* [2], which aims to identify the most relevant datasets from a data lake (a repository of data stored in its raw format). *Data integration* is another task that combines multiple datasets into a larger dataset [35]. As more and more datasets become available, we anticipate the popularity of amassing datasets of interest will grow, and so does the need for more scalable algorithms.

**Data labeling** aims to add informative labels to data samples. Manual methods, such as *crowdsourcing*, are accurate but costly. *Semi-supervised labeling* is a family of techniques that infer the labels of unlabeled data based on a small amount of labeled data [42, 15] (e.g., training a model to make predictions on the unlabeled data). *Active learning* is an iterative labeling strategy that selects the most informative unlabeled samples in each iteration [28, 43]. Other research has redefined the labeling procedure in weakly-supervised settings [27, 44]. For example, *data programming* [27], takes domain-specific heuristics as input to infer labels. Large labeled datasets are key enablers of deep learning. We predict that more efficient labeling methods with various formats of human involvement for diverse data types will be proposed.

**Data preparation** cleans and transforms raw data into a form suitable for learning. *Data cleaning* helps remove noises and errors in data [3], such as imputing missing values, dropping duplicates, and fixing inconsis-

tencies. *Feature extraction* aims to create numerical features usable for learning from raw data [30]. The extraction strategies often depend on data formats (e.g., tf-idf for text and patches for images). *Data pre-processing* techniques, such as standardization and normalization, make data more suitable for model training. Our anticipation for data preparation is that future research will focus more on automatically identifying the best strategies to deal with arbitrary data.

**Data reduction** reduces the data size, making it simpler and more understandable with potentially improved performance. From the feature perspective, common strategies include selecting a subset of features (*feature selection* [19]), transforming data to a lower-dimension space (*dimension reduction* [9]), etc. From the instance perspective, *instance selection* is a notable research topic that aims to select a subset of the data that maintains the underlying distribution [24]. As modern datasets are becoming increasingly larger (both feature and instance sizes), data reduction plays a critical role in performance and capacity optimization.

**Data augmentation** is a strategy to increase data diversity by creating modified samples without collecting more data. The existing methods can generally be divided into *basic manipulation* and *deep learning* approaches [34]. The former directly makes minor changes to the original data samples, e.g., flipping images. The latter trains deep learning models to synthesize data, e.g., training generative models to capture data distributions and generate new samples. Data augmentation brings many benefits, such as improved accuracy, increased generalization capabilities, and robustness.

We recognize that more research efforts have been paid to data processing (data reduction and augmentation) than data creation (data collection and labeling). This could be partly explained by our focus on model-centric AI in the past since processing data serves as a pre-processing step of model training. Looking forward, we anticipate that more studies on data creation will appear since it fundamentally determines data quality.

### 3 Evaluation Data Development

Assessing and understanding the model quality is crucial before deployment. In the model-centric paradigm, evaluation data typically refers to test and validation sets where some metrics (e.g., accuracy) are obtained to measure prediction performances. In the DCAI paradigm, we target a more comprehensive model evaluation with more granular insights of model capabilities beyond performance metrics. We divide our discussion into *in-distribution* and *out-of-distribution* data.

**In-distribution** evaluation data refers to testing samples that follow the same distribution as the training data. Some techniques have been proposed to create highly specific in-distribution sets for various testing purposes. One representative technique is *data slicing* [4, 29], which reserves a small part of original data based on populations (e.g., genders or races). This can provide important information about where models underperform (e.g., a model could be less accurate for females than males). Another technique is *data synthesis* [41, 6], which fabricates data following the original distribution. It is usually employed when the collected data is limited or dealing with intensive testing demands. *Algorithmic recourse* [38, 16] is an emerging task to understand the decision boundaries of models. It aims to find a hypothetical in-distribution set close to the model boundary but with different predictions. For instance, if an individual is denied a loan, algorithmic recourse seeks a close sample (e.g., with an increased account balance) that can be accepted. It helps explain predictions and prevent ethical issues and fatal errors.

**Out-of-distribution** evaluation data means the testing samples follow a distribution that differs from the training data. Numerous techniques are employed for various evaluation purposes. To assess model transferability, *data grafting* (or data fusion) [32] creates a new testing set by fusing data from similar domains that have comparable data formats. For example, we can test the transferability of a sentiment classifier with the fused data collected from the sources unseen in training [5]. Besides, to test robustness, a typical approach is *adversarial perturbation*, which aims to design samples similar to the original ones but can mislead the model [11]. The adversarial sets can reveal the high-risk weakness, provide informative features to refine the model training, and evaluate ethical issues. Furthermore, we can purposely develop data to evaluate the adaptation on *distribution shift* [36], a common issue for model deployment. This is often achieved by simulating the distribution shift through weighted sampling.

Unlike training data construction, research on evaluation data development is relatively open-ended. The strategies are often purpose-driven, aiming to test a spe-

cific property of the model, such as transferability, robustness, etc. We predict that more systematic research will be conducted in this direction and more evaluation data creation algorithms will be proposed.

### 4 Data Maintenance

In production scenarios, data is not created once, but rather continuously updated. Data maintenance is a significant challenge that DCAI has to consider. We need to organize the data and ensure its reliability in a dynamic environment. Our discussion involves data organization, data understanding, and data quality.

**Data organization** provides a reliable data source for continuous development. This is often achieved via *data warehousing* [1], which combines information from multiple sources into one comprehensive warehousing system. One crucial task is *metadata management* [37], where metadata means the context where the data is needed. It is a communication protocol with many benefits, including software reuse, information extraction, data summarization, etc. Another noteworthy task is *data version control* [17], which records the changes to a dataset during its lifecycle to achieve traceable error reporting and enable reproducible results.

**Data understanding** aims to develop algorithms and tools to help comprehensively understand and debug data. One important measure is *data visualization* [10], which represents data in a more intuitive form, such as graphics, to facilitate communication. A notable example is warehouse systems, where entity-relation diagrams have been exploited to show the relationship between data schema, and graphical representations are developed to visualize logical structure among queries and outputs. *Data interpretation* [31, 40] is another measure that provides a deeper understanding of data. It reviews data and draws conclusions from data using analytic methods, which are often domain-specific.

**Data quality** is the key to model training. *Quality assessment* aims to develop evaluation metrics to measure data quality and identify potential flaws and risks [25]. *Quality improvement* influences different stages of a data pipeline [20]. The improvement methods can be performed by domain experts (e.g., auditing and feedback), collective intelligence (e.g., majority voting), user-defined rules (e.g., data unit-test), or automated algorithms (e.g., correcting errors in labels).

Data maintenance is not an isolated component in DCAI but rather plays a fundamental and supporting role in the data ecosystem. With this perspective in mind, we predict that data maintenance strategies will become more dependent on training and evaluation data construction, providing them with continuous data support in an ever-changing environment.

## 5 Open Research Challenges

To achieve DCAI, some open research challenges still deserve future exploration in our community.

**Evaluation Data & Data Maintenance.** The majority of previous research has focused on training data development. The model-centric AI paradigm could partly drive this since many DCAI tasks were treated as pre-processing steps of model training. We argue that the role of data expands well beyond training, and evaluation data and data maintenance are equally important. These two comparatively underexplored directions are challenging since they can be open-ended and not as well defined as constructing training data. Rather than desperately optimizing performance metrics, they aim to provide a comprehensive understanding of the performance and continuous support of data.

**Cross-task Techniques.** Despite the progresses on various individual tasks, the investigation from the broader DCAI view is relatively lacking. In particular, different DCAI tasks could have an interaction effect. The optimal data augmentation choice may depend on the collected data; the evaluation set construction strategy needs to consider training data, and the training data could be adjusted based on the evaluation results; data maintenance strategies must be designed based on the training/evaluation data characteristics. It remains a challenge to systematically and simultaneously tackle data issues across multiple DCAI tasks. AutoML could be one of the promising directions to approach this goal with end-to-end data pipeline search [8, 12, 7, 18, 45].

**Data-model Co-design.** While DCAI advocates shifting our focus to data, it does not necessarily imply that the model has to stay unchanged. The optimal data strategies could differ when using different models and vice versa. With this in mind, our prediction is that future advancements will come from co-designing data pipelines and models and that the data-model co-evolution will pave the way to more powerful AI systems. More studies are encouraged to investigate data-model relationships and co-design techniques.

**Data Bias.** Recently, AI systems in many high-stake applications have been reported to show discrimination against certain groups of people, which raises serious fairness concerns [5, 39, 22]. The root cause often lies in the biased distributions of specific sensitive variables in data. From the DCAI perspective, some intriguing challenges arise: 1) How to mitigate the bias in the training data? 2) How to construct evaluation data to expose the unfairness issue? 3) How to continuously maintain data unbiasedness in a dynamic environment?

**Benchmarks.** In the model-centric paradigm, benchmarks propel us forward in advancing model designs. However, benchmarks are lacking for DCAI. The

existing benchmarks often only focus on a specific DCAI task (e.g., feature selection [19]). It remains a challenge to construct a *data benchmark* to understand the overall data quality and comprehensively evaluate various DCAI techniques [21]. Efforts in DCAI benchmarking will significantly accelerate our research progress.

## 6 Conclusion

DCAI is an emerging research field that fundamentally shifts our focus from model to data. We provide a top-level discussion of its missions, aiming to help the community understand DCAI and push forward its progress. Many unsolved challenges remain to be addressed before we can fully achieve DCAI.

## References

- [1] M. ARMBRUST, A. GHODSI, R. XIN, AND M. ZAHARIA, *Lakehouse: a new generation of open platforms that unify data warehousing and advanced analytics*, in Proceedings of CIDR, 2021.
- [2] A. BOGATU, A. A. FERNANDES, N. W. PATON, AND N. KONSTANTINOY, *Dataset discovery in data lakes*, in ICDE, 2020.
- [3] X. CHU, I. F. ILYAS, S. KRISHNAN, AND J. WANG, *Data cleaning: Overview and emerging challenges*, in SIGMOD, 2016.
- [4] Y. CHUNG, T. KRASKA, N. POLYZOTIS, K. H. TAE, AND S. E. WHANG, *Slice finder: Automated data slicing for model validation*, in ICDE, 2019.
- [5] M. DU, N. LIU, F. YANG, AND X. HU, *Learning credible deep neural networks with rationale regularization*, in ICDM, 2019.
- [6] K. EL EMAM, L. MOSQUERA, J. BASS, ET AL., *Evaluating identity disclosure risk in fully synthetic health data: model development and validation*, Journal of medical Internet research, 22 (2020), p. e23139.
- [7] M. FEURER, K. EGGENSERGER, S. FALKNER, M. LINDAUER, AND F. HUTTER, *Auto-sklearn 2.0: Hands-free automl via meta-learning*, arXiv preprint arXiv:2007.04074, (2020).
- [8] M. FEURER, A. KLEIN, K. EGGENSERGER, J. SPRINGENBERG, M. BLUM, AND F. HUTTER, *Efficient and robust automated machine learning*, NeurIPS, (2015).
- [9] I. K. FODOR, *A survey of dimension reduction techniques*, tech. report, Lawrence Livermore National Lab., CA (US), 2002.
- [10] S. GATHANI, P. LIM, AND L. BATTLE, *Debugging database queries: A survey of tools, techniques, and users*, in CHI, 2020.
- [11] I. J. GOODFELLOW, J. SHLENS, AND C. SZEGEDY, *Explaining and harnessing adversarial examples*, arXiv preprint arXiv:1412.6572, (2014).
- [12] Y. HEFFETZ, R. VAINSHTAIN, G. KATZ, AND L. ROKACH, *Deepline: Automl tool for pipelines gener-*

- ation using deep reinforcement learning and hierarchical actions filtering, in KDD, 2020.
- [13] J. JAKUBIK, M. VÖSSING, N. KÜHL, J. WALK, AND G. SATZGER, *Data-centric artificial intelligence*, arXiv preprint arXiv:2212.11854, (2022).
  - [14] M. H. JARRAHI, A. MEMARIANI, AND S. GUHA, *The principles of data-centric ai (dcai)*, arXiv preprint arXiv:2211.14611, (2022).
  - [15] G. KARAMANOLAKIS, S. MUKHERJEE, G. ZHENG, AND A. HASSAN, *Self-training with weak supervision*, in NAACL, 2021.
  - [16] A.-H. KARIMI, B. SCHÖLKOPF, AND I. VALERA, *Algorithmic recourse: from counterfactual explanations to interventions*, in FAccT, 2021.
  - [17] J. KLUMP, L. WYBORN, M. WU, J. MARTIN, R. R. DOWNS, AND A. ASMI, *Versioning data is about more than revisions: A conceptual framework and proposed principles*, Data Science Journal, (2021).
  - [18] K.-H. LAI, D. ZHA, G. WANG, J. XU, Y. ZHAO, D. KUMAR, Y. CHEN, P. ZUMKHAWAKA, M. WAN, D. MARTINEZ, ET AL., *Tods: An automated time series outlier detection system*, in AAI, 2021.
  - [19] J. LI, K. CHENG, S. WANG, F. MORSTATTER, R. P. TREVINO, J. TANG, AND H. LIU, *Feature selection: A data perspective*, ACM computing surveys (CSUR), 50 (2017), pp. 1–45.
  - [20] D. LOSHIN, *The practitioner’s guide to data quality improvement*, Elsevier, 2010.
  - [21] M. MAZUMDER, C. BANBURY, X. YAO, B. KARLAŠ, W. G. ROJAS, S. DIAMOS, G. DIAMOS, L. HE, D. KIELA, D. JURADO, ET AL., *Dataperf: Benchmarks for data-centric ai development*, arXiv preprint arXiv:2207.10062, (2022).
  - [22] N. MEHRABI, F. MORSTATTER, N. SAXENA, K. LERMAN, AND A. GALSTYAN, *A survey on bias and fairness in machine learning*, ACM Computing Surveys (CSUR), 54 (2021), pp. 1–35.
  - [23] A. NG, D. LAIRD, AND L. HE, *Data-centric ai competition*, DeepLearning AI. Available online: <https://https-deeplearning-ai.github.io/data-centric-comp/> (accessed on 9 December 2021), (2021).
  - [24] J. A. OLVERA-LÓPEZ, J. A. CARRASCO-OCHOA, J. MARTÍNEZ-TRINIDAD, AND J. KITTLER, *A review of instance selection methods*, Artificial Intelligence Review, 34 (2010), pp. 133–143.
  - [25] L. L. PIPINO, Y. W. LEE, AND R. Y. WANG, *Data quality assessment*, Communications of the ACM, 45 (2002), pp. 211–218.
  - [26] N. POLYZOTIS AND M. ZAHARIA, *What can data-centric ai learn from data and ml engineering?*, arXiv preprint arXiv:2112.06439, (2021).
  - [27] A. J. RATNER, C. M. DE SA, S. WU, D. SELSAM, AND C. RÉ, *Data programming: Creating large training sets, quickly*, in NeurIPS, 2016.
  - [28] P. REN, Y. XIAO, X. CHANG, P.-Y. HUANG, Z. LI, B. B. GUPTA, X. CHEN, AND X. WANG, *A survey of deep active learning*, ACM computing surveys (CSUR), 54 (2021), pp. 1–40.
  - [29] S. SAGADEEVA AND M. BOEHM, *Sliceline: Fast, linear-algebra-based slice finding for ml model debugging*, in SIGMOD, 2021.
  - [30] A. O. SALAU AND S. JAIN, *Feature extraction: a survey of the types, techniques, applications*, in ICSC, 2019.
  - [31] H. SALEHI, S. BISWAS, AND R. BURGUEÑO, *Data interpretation framework integrating machine learning and pattern recognition for self-powered data-driven damage identification with harvested energy variations*, Engineering Applications of Artificial Intelligence, 86 (2019), pp. 136–153.
  - [32] G. SAPORTA, *Data fusion and data grafting*, Computational statistics & data analysis, 38 (2002), pp. 465–473.
  - [33] B. SETTLES, *Active learning literature survey*, (2009).
  - [34] C. SHORTEN AND T. M. KHOSHGOFTAAR, *A survey on image data augmentation for deep learning*, Journal of big data, 6 (2019).
  - [35] M. STONEBRAKER, I. F. ILYAS, ET AL., *Data integration: The current status and the way forward.*, IEEE Data Eng. Bull., 41 (2018), pp. 3–9.
  - [36] M. SUGIYAMA, M. KRAULEDAT, AND K.-R. MÜLLER, *Covariate shift adaptation by importance weighted cross validation.*, Journal of Machine Learning Research, 8 (2007).
  - [37] A. VADUVA AND T. VETTERLI, *Metadata management for data warehousing: An overview*, International Journal of Cooperative Information Systems, 10 (2001), pp. 273–298.
  - [38] S. VENKATASUBRAMANIAN AND M. ALFANO, *The philosophical basis of algorithmic recourse*, in FAccT, 2020.
  - [39] M. WAN, D. ZHA, N. LIU, AND N. ZOU, *In-processing modeling techniques for machine learning fairness: A survey*, ACM Transactions on Knowledge Discovery from Data (TKDD).
  - [40] T. WU, E. HU, S. XU, M. CHEN, P. GUO, Z. DAI, T. FENG, L. ZHOU, W. TANG, L. ZHAN, ET AL., *clusterprofiler 4.0: A universal enrichment tool for interpreting omics data*, The Innovation, 2 (2021), p. 100141.
  - [41] L. XU, M. SKOULARIDOU, A. CUESTA-INFANTE, AND K. VEERAMACHANENI, *Modeling tabular data using conditional gan*, in NeurIPS, 2019.
  - [42] Y. XU, J. DING, L. ZHANG, AND S. ZHOU, *Dp-ssl: Towards robust semi-supervised learning with a few labeled samples*, in NeurIPS, 2021.
  - [43] D. ZHA, K.-H. LAI, M. WAN, AND X. HU, *Metaaad: Active anomaly detection with deep reinforcement learning*, in ICDM, 2020.
  - [44] D. ZHA AND C. LI, *Multi-label dataless text classification with topic modeling*, Knowledge and Information Systems, 61 (2019), pp. 137–160.
  - [45] D. ZHA, Z. PERVAIZ BHAT, Y.-W. CHEN, Y. WANG, S. DING, A. K. JAIN, M. QAZIM BHAT, K.-H. LAI, J. CHEN, ET AL., *Autovideo: An automated video action recognition system*, in IJCAI, 2022.